

AI Worklog

Shyamraj H · NxtWave Growth Intern challenge, round 1 · Plan: 500 final-years for a free AI workshop in 7 days on ₹2,000

Tools: Claude (Claude Code) for research, planning, building and testing; Google Sheets + Apps Script for the backend; headless Chrome for screenshots; PowerPoint for the deck. Every number in the plan comes from a simulator inside the asset, and the code is covered by 10 automated tests.

Three examples: prompt → AI output → what I changed → final output

1. Choosing what to build

Prompt	Here is the NxtWave growth challenge brief. Finish the whole task and ask me what you need to know.
AI output	Four asset options (referral engine + dashboard, WhatsApp bot flow, live workshop engagement tool, AI project evaluator), recommending the referral engine because the brief warns that everyone will build a landing page.
What I changed / rejected	Chose the referral engine. Rejected the WhatsApp bot (hard to show as a live, working link) and the AI evaluator (impressive, but it doesn't move registrations). Kept the landing page only as the engine's front door.
Final output	A sign-up page with personal referral links and rewards, plus a growth console that runs the plan's kill/scale rules as "Today's decisions".

2. Where the data lives

Prompt	Make the live link store real sign-ups, so a reviewer can register and see it work.
AI output	First, a Claude-hosted artifact database (zero setup). It dropped that itself after checking that only signed-in claude.ai users can write to it, so students couldn't register. Then two options: Supabase (real Postgres, live leaderboard) or Google Sheets + Apps Script.
What I changed / rejected	Rejected Supabase. Free projects pause after about a week without traffic, so a reviewer opening my link later could hit a dead backend. Chose Google Sheets: free, never sleeps, and the growth team can open the sheet without learning a new tool.
Final output	An Apps Script web app on a Google Sheet (Registrations, Visits and Partners tabs) with duplicate blocking, fraud flags, privacy-safe public data and an admin key for the console.

3. Stress-testing the plan's math

Prompt	Simulate the 7-day campaign from the plan's assumptions and tell me whether it really reaches 500.
AI output	A simulator and a plan landing at about 515, a 3% buffer, plus a rule that flagged any channel where under 20% of visitors signed up.
What changed	Review of the first run caught two problems. Champion sign-ups came out 9% below the plan's own arithmetic: a sampling bug, now fixed. The conversion rule flagged the paid posts every single day, but cold traffic always converts lower, so paid is now judged on ₹ per sign-up.
Final output	100 simulated campaigns: median 548, 90 clear 500, worst 463. The deck states those odds and a downside plan, and the console ships a plan simulator so anyone can change the assumptions.

How I thought

What changed between the first idea and the final solution?

- **From landing page to growth system.** The first version was a sign-up page. The final version is a referral engine plus a console that runs the plan: every partner gets a tracked link, and the console says each night what to kill, scale or nudge.
- **From "about 515" to "about 550, with odds".** Simulating the plan 100 times showed it misses 500 in 10% of runs. So the plan states its odds and carries backup levers (5 backup champions, a second WhatsApp wave) with a trigger: day 4 more than 10% behind.
- **From "spend the budget" to "pay for outcomes".** Half the ₹2,000 became prizes, paid only after the workshop and counting only friends who attend. Paid reach became a ₹800 test with a kill rule, not a pillar.

If I had another 24 hours, what would I improve?

- **Run a one-day pilot.** Three champions in my own college would replace the assumed click-through and conversion rates with measured ones before spending a rupee.
- **Automate the WhatsApp reminders.** Trigger the day-before, 2-hour and 15-minute reminders from the Google Sheet through the WhatsApp Business API (or n8n). Today they are a manual step.
- **Verify phone numbers with OTP.** This would make referral credit fraud-proof instead of flag-and-review.
- **Close the loop on attendance.** Import the attendance list so prize eligibility and show-up rate by channel compute themselves.

What did AI suggest that I rejected, and why?

Supabase as the backend (exchange 2). It is the "proper" database choice, but free projects pause when idle, and this link has to work whenever a reviewer opens it. A Google Sheet never sleeps, and the growth team can read it on day one.

Also caught in review: judging paid traffic on landing-page conversion (exchange 3). Cold audiences always convert lower than a friend's forward, so that rule would have killed a channel that was actually cheap per sign-up. Paid is judged on cost per sign-up instead.

Where I chose not to use AI

The nightly decisions (kill a paid post, replace a silent champion, hold a suspicious referrer's reward) are fixed rules written in advance, not model output. They decide where money and messages go, so they need to be predictable, explainable and the same every night. AI helped build and stress-test the system; it does not make the calls inside it.